

Go Data Concurrency Detection

Deadlock and Race Detectors

25 April 2020

Emil Sharifullin

Lead Software Engineer, SKB Kontur

Go Go Go

- Easy to learn
- Have brilliant concurrency paradigm
- If you have more than one core concurrency become parallelism

Too many circumstances to create a bunch of concurrency errors => **Go needs tools to prevent it**

Deadlocks

Deadlocks

In concurrent computing, a deadlock is a state in which each member of a group is waiting for another member, including itself, to take action, such as sending a message or more commonly releasing a lock

```
func main() {  
    ch1 := make(chan int)  
    ch2 := make(chan int)  
    go func() {  
        fmt.Println(<-ch1)  
        ch2 <- 1  
    }()  
    fmt.Println(<-ch2)  
    ch1 <- 1  
}
```

Run

Deadlocks

Deadlocks can appear even in single threaded environment

```
func main() {  
    ch := make(chan int)  
    ch <- 1  
    fmt.Println(<-ch)  
}
```

Run

```
func main() {  
    var m1 sync.Mutex  
    m1.Lock()  
    m1.Lock()  
}
```

Run

Scheduler

Go scheduler operates with P's, M's and G's

- **G**: goroutine
- **M**: OS thread (machine) can execute at most one G at a time
- **P**: logical processor can hold at most one M at a time

At any time, there are at most GOMAXPROCS number of P

Deadlock Detector

Go runtime have a deadlock detector implemented in goroutines scheduler

```
4430 // Check for deadlock situation.  
4431 // The check is based on number of running M's, if 0 -> deadlock.  
4432 // sched.lock must be held.  
4433 func checkdead() {
```

Go Runtime <https://github.com/golang/go/blob/master/src/runtime/proc.go>

(<https://github.com/golang/go/blob/885099d1550dad8387013c8f35ad3d4ad9f17c66/src/runtime/proc.go#L4433-L4531>)

Counting M's

Running M's is a basically difference between all M's and idled M's

```
4463      run := mcount() - sched.nmiddle - sched.nmiddlelocked - sched.nmsys
4464      if run > run0 {
4465          return
4466      }
```

And check for system purpose M

```
4449      // If we are not running under cgo, but we have an extra M then account
4450      // for it. (It is possible to have an extra M on Windows without cgo to
4451      // accommodate callbacks created by syscall.NewCallback. See issue #6751
4452      // for details.)
4453      var run0 int32
```

If any M is running -> **No deadlock**

Checking G's

Is it any G alive in runtime?

```
4474     for i := 0; i < len(allgs); i++ {
4475         gp := allgs[i]
4476         if isSystemGoroutine(gp, false) {
4477             continue
4478         }
4479         s := readgstatus(gp)
4480         switch s &^ _Gscan {
4481             case _Gwaiting,
4482                 _Gpreempted:
4483                 grunning++
```

If no G is running -> **Deadlock**

```
4493     if grunning == 0 { // possible if main goroutine calls runtime.Goexit()
4494         unlock(&sched.lock) // unlock so that GODEBUG=scheddetail=1 doesn't hang
4495         throw("no goroutines (main called runtime.Goexit) - deadlock!")
4496     }
```

Checking P's

Eventually check that any P have actions to take at some time

```
4522     for _, _p_ := range allp {  
4523         if len(_p_.timers) > 0 {  
4524             return  
4525         }  
4526     }
```

If some action scheduled -> **No deadlock**

Deadlock

If not returned before -> **Deadlock!**

```
4528     getg().m.throwing = -1 // do not dump full stacks
4529     unlock(&sched.lock)    // unlock so that GODEBUG=scheddetail=1 doesn't hang
4530     throw("all goroutines are asleep - deadlock!")
```

Additional Cases

Do not check if panicking

```
4445     if panicking > 0 {  
4446         return  
4447     }
```

Do not check if code builded as .so library or archive

```
4437     if islibrary || isarchive {  
4438         return  
4439     }
```

And additional P's timers case for code running in playground

```
4498     // Maybe jump time forward for playground.  
4499     if faketime != 0 {
```

Not a Magic

```
func main() {  
    ch1 := make(chan int)  
    go func() {  
        fmt.Println(<-ch1)  
    }()  
    sigs := make(chan os.Signal, 1)  
    signal.Notify(sigs, syscall.SIGINT, syscall.SIGTERM)  
    <-sigs  
}
```

Run

Data Races

Data Races

A data race occurs when two threads access the same variable concurrently and at least one of the accesses is write.

```
counter := 0
var wg sync.WaitGroup
for i := 0; i < 10000; i++ {
    wg.Add(1)
    go func() {
        counter = counter + 1
        wg.Done()
    }()
}
wg.Wait()
fmt.Println(counter)
```

Run

Data Race Detector

Let's check it

```
#!/usr/bin/env bash
```

```
go run -race src/race/first/main.go
```

Run

It's a result of ThreadSanitizer algorithm

ThreadSanitizer – data race detection in practice by [Serebryany & Iskhodzhanov](#)

(<https://static.googleusercontent.com/media/research.google.com/en//pubs/archive/35604.pdf>)

Happens Before

Is a way to order events between different executors(processes/threads/goroutines) in concurrent(distributed) systems

Event **E1** happens before event **E2**($E1 < E2$) if at least one of following is true

- **E1** and **E2** in one executor and **E1** placed earlier than **E2**
- **E1** and **E2** is a Lock-Unlock events of one synchronization primitive
- Exist **E'** that $E1 < E'$ and $E' < E2$

In the other cases **E1** and **E2** are concurrent



Vector Clocks

A vector clock is an algorithm for generating a partial ordering of events in a distributed system and detecting causality violations

It's based on Lamport clock idea and was developed by Colin Fidge and Friedemann Mattern

Allow us to determine Happens Before relation in distributed systems

Vector Clocks in Practice



So $E1 < E2$ if $V[E1] < V[E2]$

$V[E1] = [1,0]$, $V[E2] = [2,1]$

$V[E1] < V[E2] \Rightarrow E1 < E2$

$V[E3] = [1,0]$, $V[E4] = [0,1]$

$V[E3] \not< V[E4] \Rightarrow E3 \not< E4$

$V[E4] \not< V[E3] \Rightarrow E4 \not< E3 \Rightarrow$ Concurrent!

Applying to Go

All these concepts are implemented in Race Detector that called ThreadSanitizer

- **E1..En** - Memory reads and writes
- **S1..Sn** - Mutex Lock/Unlock & Goroutines creation

On E1..En we compare current event with previous events that affect this memory location and do following checks:

- Is at least one of events **write**?
- Can we identify **happens before** relation?

If answers **Yes && No** => **Data Race**

Implementation

182	0x001d 00029 (src/race/first/main.go:15)	PCDATA	\$0, \$1
183	0x001d 00029 (src/race/first/main.go:15)	PCDATA	\$1, \$1
184	0x001d 00029 (src/race/first/main.go:15)	MOVQ	"". &counter+32(SP), AX
185	0x0022 00034 (src/race/first/main.go:15)	PCDATA	\$0, \$0
186	0x0022 00034 (src/race/first/main.go:15)	INCQ	(AX)

297	0x0026 00038 (src/race/first/main.go:14)	CALL	runtime.racefuncenter(SB)
298	0x002b 00043 (src/race/first/main.go:15)	PCDATA	\$0, \$1
299	0x002b 00043 (src/race/first/main.go:15)	MOVQ	"". &counter+40(SP), AX
300	0x0030 00048 (src/race/first/main.go:15)	PCDATA	\$0, \$0
301	0x0030 00048 (src/race/first/main.go:15)	MOVQ	AX, (SP)
302	0x0034 00052 (src/race/first/main.go:15)	CALL	runtime.raceread(SB)
303	0x0039 00057 (src/race/first/main.go:15)	PCDATA	\$0, \$1
304	0x0039 00057 (src/race/first/main.go:15)	MOVQ	"". &counter+40(SP), AX
305	0x003e 00062 (src/race/first/main.go:15)	MOVQ	(AX), CX
306	0x0041 00065 (src/race/first/main.go:15)	MOVQ	CX, ""..autotmp_7+16(SP)
307	0x0046 00070 (src/race/first/main.go:15)	PCDATA	\$0, \$0
308	0x0046 00070 (src/race/first/main.go:15)	MOVQ	AX, (SP)
309	0x004a 00074 (src/race/first/main.go:15)	CALL	runtime.racewrite(SB)
310	0x004f 00079 (src/race/first/main.go:15)	MOVQ	""..autotmp_7+16(SP), AX
311	0x0054 00084 (src/race/first/main.go:15)	INCQ	AX
312	0x0057 00087 (src/race/first/main.go:15)	PCDATA	\$0, \$2
313	0x0057 00087 (src/race/first/main.go:15)	PCDATA	\$1, \$1
314	0x0057 00087 (src/race/first/main.go:15)	MOVQ	"". &counter+40(SP), CX

Implementation

Notify race detector on a new goroutine creation and mutex Lock/Unlock

```
3563     if raceenabled {  
3564         newg.racectx = racegostart(callerpc)  
3565     }
```

func newproc1 <https://github.com/golang/go/blob/master/src/runtime/proc.go>

(<https://github.com/golang/go/blob/885099d1550dad8387013c8f35ad3d4ad9f17c66/src/runtime/proc.go#L3563-L3565>)

```
75     if race.Enabled {  
76         race.Acquire(unsafe.Pointer(m))  
77     }
```

func (m *Mutex) Lock <https://github.com/golang/go/blob/master/src/sync/mutex.go>

(<https://github.com/golang/go/blob/885099d1550dad8387013c8f35ad3d4ad9f17c66/src/sync/mutex.go#L75-L77>)

```
180     if race.Enabled {  
181         _ = m.state  
182         race.Release(unsafe.Pointer(m))  
183     }
```

func (m *Mutex) Unlock <https://github.com/golang/go/blob/master/src/sync/mutex.go>

(<https://github.com/golang/go/blob/885099d1550dad8387013c8f35ad3d4ad9f17c66/src/sync/mutex.go#L180-L183>)

Fixed Example

```
counter := 0
var wg sync.WaitGroup
var m sync.Mutex

for i := 0; i < 10000; i++ {
    wg.Add(1)
    go func() {
        m.Lock()
        counter = counter + 1
        m.Unlock()
        wg.Done()
    }()
}
wg.Wait()
fmt.Println(counter)
```

Run

```
#!/usr/bin/env bash
```

```
go run -race src/race/second/main.go
```

Run

Pros and Cons

- Do not have false positives
 - Widely used and do not have obvious bugs
-
- Some races can be not found
 - Memory consumption increases 5-10 times
 - Execution time increases 5-15 times
 - ThreadSanitizer is a C++ library so works only with CGo enabled

When and How to Use

Deadlock detector

- Anytime
- Anywhere

Data race detector

- In tests **go test -race ./...**
- During development **go run -race main.go**
- On special environment for race detection **go build -race .**

Thank you

Deadlock and Race Detectors

25 April 2020

Emil Sharifullin

Lead Software Engineer, SKB Kontur

iam@litleleprikon.me (mailto:iam@litleleprikon.me)

[@litleleprikon](http://twitter.com/litleleprikon) (http://twitter.com/litleleprikon)

